

Experiment 4

k-Nearest Neighbors (KNN)

Aim:

To implement the k-Nearest Neighbours (KNN) algorithm for classifying the Iris dataset and evaluate its performance using standard classification metrics.

Theory:

The k-Nearest Neighbours (KNN) algorithm is a supervised machine learning algorithm used for classification and regression. It is a non-parametric method, as it does not assume any prior distribution of the data. In KNN, classification is performed based on the similarity between data samples, which is commonly measured using distance metrics such as Euclidean distance.

The nearest neighbour classification concept was introduced by Fix and Hodges and later extended by Cover and Hart. In this method, a test sample is assigned a class label by identifying the k nearest training samples in the feature space and applying majority voting. The value of k represents the number of nearest neighbours considered for classification.


KNN does not involve an explicit training phase. Instead, it stores the entire training dataset and performs computation during the testing phase. Hence, it is referred to as a lazy learning or instance based learning algorithm. The choice of k significantly affects classifier performance. Smaller values of k may result in overfitting, while larger values of k may lead to underfitting.

Since KNN is based on distance calculations, feature scaling is important to ensure that all features contribute equally. Without proper scaling, features with larger numerical ranges may dominate the distance computation.

Table: Merits and Demerits of k-Nearest Neighbours (KNN)

Merits of KNN	Demerits of KNN
Simple and easy to understand	High computational cost during prediction
Does not require an explicit training phase	Requires large memory to store training data
Makes no assumption about data distribution	Highly sensitive to feature scaling
Works well for small and well-separated datasets	Performance degrades in high-dimensional data
Can be used for both classification and regression	Sensitive to noise and outliers

Explanation of the figure: The figure illustrates how K-NN assigns a class to a new data point by considering the majority class among its nearest neighbours.



Code Section:

1. INPUT: Import Libraries

```
# Import Libraries

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

print ("Imported Analysis and Plotting libraries")


from sklearn.datasets import load_iris

from sklearn.preprocessing import StandardScaler

from sklearn.preprocessing import label_binarize
```

```
from sklearn.model_selection import train_test_split

from sklearn.model_selection import StratifiedKFold

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import (

    accuracy_score,

    precision_score,

    recall_score,

    f1_score,

    roc_auc_score,

    classification_report,

    confusion_matrix

)

print ("Imported Sklearn Modules")
```

OUTPUT:

Imported Analysis and Plotting libraries

Imported Sklearn Modules

INPUT: Loading and Reading the Dataset

```
# Loading and Reading Data
```

```
iris = load_iris () # returns data in numpy arrays
```

```
X = iris.data
```

```
y = iris.target
```

```
print ("Extracted Feature and Target Arrays")
```

```
names = iris.feature_names # extracting feature names
```

```
target_names = iris.target_names # extracting target names
```

```
iris_df = pd.DataFrame(X, columns=names) # making a dataframe with column names
```

```
iris_df["label"] = [target_names[i] for i in y]
```

```
iris_df.head()
```

OUTPUT:

Extracted Feature and Target as numpy Arrays

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	label
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

INPUT: Exploratory Data Analysis

```
# Data Analysis
```

```
iris_df.info () # Data Overview
```

```
iris_df["label"].value_counts () # Target counts
```

```
iris_df.describe() # Dataset description
```

```
iris_df.isnull(). sum () # Finding any null values
```

OUTPUT:

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 150 entries, 0 to 149
```

```
Data columns (total 5 columns):
```

```
# Column Non-Null Count Dtype
```

```
-----
```

```
0 sepal length (cm) 150 non-null float64
```

```
1 sepal width (cm) 150 non-null float64
```

```
2 petal length (cm) 150 non-null float64
```

```
3 petal width (cm) 150 non-null float64
```

```
4 label 150 non-null object
```

```
dtypes: float64(4), object(1)
```

```
memory usage: 6.0+ KB
```

```

label
setosa                                     50
versicolor                               50
virginica                                 50
Name: count, dtype: int64

```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```

sepal          length          (cm)          0
sepal          width           (cm)          0
petal          length          (cm)          0
petal          width           (cm)          0
label                                                  0
dtype: int64

```

INPUT: Feature Distribution Visualization

```
fig, axes = plt.subplots(2, 2, figsize= (10, 8))
```

```
for ax, col in zip(axes.ravel(), names):
```

```
    sns.histplot(
```

```
        data=iris_df, x=col, hue="label",
```

```
        kde=True, element="step", ax=ax
```

```
    )
```

```
    ax.set_xlabel (col, fontweight='bold')
```


```
    ax.set_ylabel ("Count", fontweight='bold')
```

```
plt.suptitle("Feature Distributions by Class", y=1.02, fontweight='bold')
```

```
plt.tight_layout()
```

```
plt.show()
```

Feature Distributions by Class



INPUT: Feature Correlation Heatmap

```
plt.figure(figsize= (10, 8))
```

```
ax = sns.heatmap(
```

```
    iris_df.iloc[: :4]. corr (),
```

```
    annot=True,
```

```
    cmap="crest",
```


```
    fmt=".2f",
```

```
    linewidths=0.7,
```

```
    annot_kws= {"color": "black"} # Force annotation text to be black
```

```
)  
  
plt.title("Feature Correlation Heatmap", fontweight='bold')  
  
plt.show()
```

OUTPUT:



INPUT: Data Encoding

```
# Data Encoding
```

```
iris_df["target"] = y # we already have encoded targets from sklearn API call
```

```
iris_df.head()
```

OUTPUT:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	label	target
0	5.1	3.5	1.4	0.2	setosa	0
1	4.9	3.0	1.4	0.2	setosa	0
2	4.7	3.2	1.3	0.2	setosa	0
3	4.6	3.1	1.5	0.2	setosa	0
4	5.0	3.6	1.4	0.2	setosa	0

INPUT: Data Preprocessing and Train–Test Split:

```
# Data Preprocessing
```

```
X = iris_df.iloc[:4]. values
```

```
y = iris_df["target"]. values
```

```
X_train, X_test, y_train, y_test = train_test_split(
```

```
    X, y, test_size=0.3, stratify=y, # Important as it decides the behaviour of model
```

```
    random_state=42
```

```
)
```

```
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

OUTPUT:

```
(105, 4) (45, 4) (105,) (45,)
```

INPUT: Feature Scaling

```
# Scaling the values

scaler = StandardScaler ()

scaler

X_train_s = scaler.fit_transform(X_train)

X_test_s = scaler.transform(X_test)

print ("Values scaled with Standard Scaler")

X_train [:5]

X_train_s [:5]
```

OUTPUT:

```
StandardScaler()
```

```
Values scaled with Standard Scaler
```

```
array([[5.1, 2.5, 3., 1.1],
 [6.2, 2.2, 4.5, 1.5],
 [5.1, 3.8, 1.5, 0.3],
 [6.8, 3.2, 5.9, 2.3],
 [5.7, 2.8, 4.1, 1.3]])
```

```
array([[-0.90045861, -1.22024754, -0.4419858 , -0.13661044],
 [ 0.38036614, -1.87955796,  0.40282929,  0.38029394],
 [-0.90045861,  1.63676428, -1.2868009 , -1.17041921],
 [ 1.07899781,  0.31814344,  1.19132338,  1.41410271],
 [-0.20182693, -0.56093712,  0.17754527,  0.12184175]])
```

INPUT: Label Binarization for ROC-AUC

```
# Binarizing encoded labels for ROC-AUC metric calculation
```

```
y_test_bin = label_binarize (y_test, classes= [0,1,2])
```

```
y_test_bin [:5]
```

OUTPUT:

```
array([[0, 0, 1],
```

```
[0, 1, 0],
```

```
[0, 0, 1],
```

```
[0, 1, 0],
```

```
[0, 0, 1]])
```

INPUT: Model Training

```
# Model Training
```

```
print (" | | Cell in development | |")
```

```
model = KNeighborsClassifier(n_neighbors=10) # k = 10
```

```
model.fit (X_train_s, y_train)
```

```
print (" | | Animation Here | |")
```

OUTPUT:

```
Model trained using Euclidean distance with k=10
```

```
KNeighborsClassifier(metric='euclidean', n_neighbors=10)
```

INPUT: Model Evaluation and Prediction

```
# Model Evaluation
```

```
y_pred = model.predict(X_test_s)
```

```
y_pred
```

```
y_prob = model.predict_proba(X_test_s)
```

```
y_prob[:5]
```

```
print ("\nClassification Report (k=10):")
```

```
print(classification_report (
```

```
    y_test,
```

```
    y_pred,
```

```
    target_names=target_names,
```

```
    digits=4
```

```
))
```

OUTPUT:

```
array([2, 1, 1, 1, 2, 2, 1, 1, 0, 2, 0, 0, 2, 2, 0, 2, 1, 0, 0, 0, 1, 0,
1, 2, 1, 1, 1, 1, 1, 0, 2, 2, 1, 0, 2, 0, 0, 0, 0, 1, 1, 0, 1, 2,
1])
```

```
array([[0. , 0.1, 0.9],
[0. , 0.9, 0.1],
[0. , 0.7, 0.3],
[0. , 0.7, 0.3],
[0. , 0.4, 0.6]])
```

Classification Report (Euclidean, k=10):

	precision	recall	f1-score	support
setosa	1.0000	1.0000	1.0000	15
versicolor	0.8333	1.0000	0.9091	15
virginica	1.0000	0.8000	0.8889	15
accuracy		0.9333		45
macro avg	0.9444	0.9333	0.9327	45
weighted avg	0.9444	0.9333	0.9327	45

INPUT: Confusion Matrix Visualization

```
plt.figure(figsize= (7, 5))

sns.heatmap(

    confusion_matrix (y_test, y_pred), annot=True, fmt="d", cmap="Greens",

    xticklabels=target_names, yticklabels=target_names

)

plt.xlabel("Predicted", fontweight='bold')


plt.ylabel("True", fontweight='bold')

plt.title("Confusion Matrix Heatmap (KNN, k=10)", fontweight='bold')

plt.tight_layout ()

plt.show()
```

OUTPUT:



13. Final Evaluation Metrics:

```
# Final Metrics and how visualization of learning
print(f'''Accuracy: {accuracy_score(y_test, y_pred)}
Precision: {precision_score(y_test, y_pred, average="macro")}
Recall: {recall_score(y_test, y_pred, average="macro")}
F1: {f1_score(y_test, y_pred, average="macro")}
ROC: {roc_auc_score(y_test_bin, y_prob, multi_class="ovr", average="macro")}
''')
```

Print ("\\n|| Cell in development ||")

Print ("|| Animation here ||")

Metric	Value
Accuracy	0.9333
Precision (Macro Avg)	0.9444
Recall (Macro Avg)	0.9333
F1-Score (Macro Avg)	0.9327
ROC-AUC (OVR)	0.9948

KNN - Practical Implementation

Virtual Labs - Dayalbagh Educational Institute

Step 1: Importing Libraries

Import Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
print("Imported Analysis and Plotting libraries")
```

Output:

```
Imported Analysis and Plotting libraries
```

Import scikit-learn modules

```
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import label_binarize
from sklearn.model_selection import train_test_split
from sklearn.model_selection import StratifiedKFold
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    roc_auc_score,
    classification_report,
    confusion_matrix
)
print("Imported Sklearn Modules")
```

Output:

```
Imported Sklearn Modules
```

Step 2: Loading Dataset

Loading and Reading Data

```
iris = load_iris()
X = iris.data
y = iris.target
print("Extracted Feature and Target as numpy Arrays")
```

Output:

```
Extracted Feature and Target as numpy Arrays
```

extracting feature names

```

names = iris.feature_names

# extracting target names
target_names = iris.target_names

# making a dataframe with column names
iris_df = pd.DataFrame(X, columns=names)
iris_df["label"] = [target_names[i] for i in y]
iris_df.head()

```

Output:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	label
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

Step 3: Data Analysis

Data Analysis

```

# Overview
iris_df.info()

```

Output:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sepal length (cm)      150 non-null   float64
1   sepal width (cm)       150 non-null   float64
2   petal length (cm)      150 non-null   float64
3   petal width (cm)       150 non-null   float64
4   label                  150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB

```

Target counts

```
iris_df["label"].value_counts()
```

Output:

```

label
setosa      50
versicolor  50
virginica   50
Name: count, dtype: int64

```

Dataset description

```
iris_df.describe()
```

Output:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

Finding any null values

```
iris_df.isnull().sum()
```


Output:

```
sepal length (cm)    0
sepal width (cm)     0
petal length (cm)    0
petal width (cm)     0
label                0
dtype: int64
```

Plotting a histogram distribution of data points with labelled classes


```
fig, axes = plt.subplots(2, 2, figsize=(10, 8))
for ax, col in zip(axes.ravel(), names):
    sns.histplot(
        data=iris_df, x=col, hue="label",
        kde=True, element="step", ax=ax
    )
    ax.set_xlabel(col, fontweight='bold')
    ax.set_ylabel("Count", fontweight='bold')
plt.suptitle("Feature Distributions by Class", y=1.02, fontweight='bold')
plt.tight_layout()
plt.show()
```

Feature Distributions by Class



Plotting correlation between features

```
plt.figure(figsize=(10, 8))
ax = sns.heatmap(
    iris_df.iloc[:, :4].corr(),
    annot=True,
    cmap="crest",
    fmt=".2f",
    linewidths=0.7,
    # Force annotation text to be black
    annot_kws={"color": "black"}
)
plt.title("Feature Correlation Heatmap", fontweight='bold')
plt.show()
```



Step 4: Data Preprocessing

Data Preprocessing

```
# Data Encoding
# we already have encoded targets from sklearn datasets
iris_df["target"] = y
iris_df.head()
```

Output:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	label	target
0	5.1	3.5	1.4	0.2	setosa	0
1	4.9	3.0	1.4	0.2	setosa	0
2	4.7	3.2	1.3	0.2	setosa	0
3	4.6	3.1	1.5	0.2	setosa	0
4	5.0	3.6	1.4	0.2	setosa	0

Stratification is important to balance classes between split

```
X = iris_df.iloc[:, :4].values
y = iris_df["target"].values

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3,
    stratify=y,
    random_state=42
)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

Output:

```
(105, 4) (45, 4) (105,) (45,)
```

Scaling the values

```
scaler = StandardScaler()
scaler
```

Output:

```
StandardScaler()
```

Apply scaling to training and test data

```
X_train_s = scaler.fit_transform(X_train)
X_test_s = scaler.transform(X_test)
print("Values scaled with Standard Scaler")
```

Output:

```
Values scaled with Standard Scaler
```

View original training data (before scaling)

```
X_train[:5]
```

Output:

```
array([[5.1, 2.5, 3. , 1.1],
       [6.2, 2.2, 4.5, 1.5],
       [5.1, 3.8, 1.5, 0.3],
       [6.8, 3.2, 5.9, 2.3],
       [5.7, 2.8, 4.1, 1.3]])
```

View scaled training data (after scaling)

```
X_train_s[:5]
```

Output:

```
array([[ -0.90045861, -1.22024754, -0.4419858 , -0.13661044],
       [ 0.38036614, -1.87955796,  0.40282929,  0.38029394],
       [-0.90045861,  1.63676428, -1.2868009 , -1.17041921],
       [ 1.07899781,  0.31814344,  1.19132338,  1.41410271],
       [-0.20182693, -0.56093712,  0.17754527,  0.12184175]])
```

Binarizing encoded labels for ROC-AUC metric calculation

```
y_test_bin = label_binarize(y_test, classes=[0,1,2])
y_test_bin[:5]
```

Output:

```
array([[0, 0, 1],
       [0, 1, 0],
       [0, 0, 1],
       [0, 1, 0],
       [0, 0, 1]])
```

Step 5: Model Training

Model Training

```
# k = 10 (Heuristic Rule)
model = KNeighborsClassifier(n_neighbors=10)
model.fit(X_train_s, y_train)
```

Output:

```
KNeighborsClassifier(n_neighbors=10)
```

Step 6: Model Evaluation

Model Evaluation

```
y_pred = model.predict(X_test_s)
y_pred
```

Output:

```
array([2, 1, 1, 1, 2, 2, 1, 1, 0, 2, 0, 0, 2, 2, 0, 2, 1, 0, 0, 0, 1, 0,
       1, 2, 1, 1, 1, 1, 1, 0, 2, 2, 1, 0, 2, 0, 0, 0, 0, 1, 1, 0, 1, 2,
       1])
```

Get prediction probabilities

```
y_prob = model.predict_proba(X_test_s)
y_prob[:5]
```

Output:

```
array([[0. , 0.1, 0.9],
       [0. , 0.9, 0.1],
       [0. , 0.7, 0.3],
       [0. , 0.7, 0.3],
       [0. , 0.4, 0.6]])
```

Generate classification report


```
print("\n\nClassification Report (k=10):")
print(classification_report(
    y_test,
    y_pred,
    target_names=target_names,
    digits=4
))
```

Output:

```
Classification Report (k=10):
precision    recall  f1-score   support
setosa       1.0000    1.0000    1.0000        15
versicolor   0.8333    1.0000    0.9091        15
virginica     1.0000    0.8000    0.8889        15
accuracy                   0.9333        45
macro avg    0.9444    0.9333    0.9327        45
weighted avg    0.9444    0.9333    0.9327        45
```

Visualize confusion matrix heatmap

```
plt.figure(figsize=(7, 5))
sns.heatmap(
    confusion_matrix(y_test, y_pred), annot=True, fmt="d", cmap="Greens",
    xticklabels=target_names, yticklabels=target_names
)
plt.xlabel("Predicted", fontweight='bold')
plt.ylabel("True", fontweight='bold')
plt.title("Confusion Matrix Heatmap (KNN, k=10)", fontweight='bold')
plt.tight_layout()
plt.show()
```



Step 7: Model Simulation

Final Metrics

```
# Accuracy will be same as F1 score as dataset is perfectly balanced and stratified
print(f'''Accuracy   : {accuracy_score(y_test, y_pred)}
Precision  : {precision_score(y_test, y_pred, average="macro")}
Recall     : {recall_score(y_test, y_pred, average="macro")}
F1         : {f1_score(y_test, y_pred, average="macro")}
ROC        : {roc_auc_score(y_test_bin, y_prob, multi_class="ovr", average="macro")}
''')
```

Output:

```
Accuracy   : 0.9333333333333333
Precision  : 0.9444444444444445
Recall     : 0.9333333333333332
F1         : 0.9326599326599326
ROC        : 0.9948148148148149
```